

実装演習 2-2. 学習率最適化手法

```
In [1]: import sys, os
sys.path.append(os.pardir) # 親ディレクトリのファイルをインポートするための設定
import numpy as np
from collections import OrderedDict
from common import layers
from data.mnist import load_mnist
import matplotlib.pyplot as plt
from multi_layer_net import MultiLayerNet
plt.style.use('ggplot')
```

```
In [2]: # MNISTデータの読み込み
(x_train, d_train), (x_test, d_test) = load_mnist(normalize=True, one_hot_label=True)
```

```
In [4]: print(f'x_train.shape: {x_train.shape}')
print(f'd_train.shape: {d_train.shape}')
print(f'x_test.shape: {x_test.shape}')
print(f'd_test.shape: {d_test.shape}')
```

```
x_train.shape: (60000, 784)
d_train.shape: (60000, 10)
x_test.shape: (10000, 784)
d_test.shape: (10000, 10)
```

オフティマイザ : SGD

```
In [5]: use_batchnorm = False

# MNISTの画像サイズは28*28=784
# 出力は0~9の分類なので10
# 活性化関数はsigmoid
network = MultiLayerNet(input_size=784, hidden_size_list=[40, 20], output_size=10, activation='sigmoid',
                        weight_init_std=0.01, use_batchnorm=use_batchnorm)
```

```
In [12]: iters_num = 1000 # 繰り返し回数
train_size = x_train.shape[0] # 学習データ (MNIST) のサイズ
batch_size = 100 # ミニバッチのサイズ
learning_rate = 0.01 # 学習率

plot_interval=10 # グラフX軸の間隔
```

```

In [15]: train_loss_list = []
        accuracies_train = []
        accuracies_test = []

        for i in range(iters_num):
            batch_mask = np.random.choice(train_size, batch_size)
            x_batch = x_train[batch_mask]
            d_batch = d_train[batch_mask]

            # 勾配
            grad = network.gradient(x_batch, d_batch)

            for key in ('W1', 'W2', 'W3', 'b1', 'b2', 'b3'):
                network.params[key] -= learning_rate * grad[key]

            loss = network.loss(x_batch, d_batch)
            train_loss_list.append(loss)

            if (i + 1) % plot_interval == 0:
                accr_test = network.accuracy(x_test, d_test)
                accuracies_test.append(accr_test)
                accr_train = network.accuracy(x_batch, d_batch)
                accuracies_train.append(accr_train)

                if (i + 1) % (plot_interval * 10) == 0:
                    print(f'Generation: {i+1}. 正答率(train)={accr_train}')
                    print(f'                : {i+1}. 正答率(test)={accr_test}')


```

```

Generation: 100. 正答率(train)=0.12
                : 100. 正答率(test)=0.1135
Generation: 200. 正答率(train)=0.12
                : 200. 正答率(test)=0.1135
Generation: 300. 正答率(train)=0.13
                : 300. 正答率(test)=0.1135
Generation: 400. 正答率(train)=0.14
                : 400. 正答率(test)=0.1135
Generation: 500. 正答率(train)=0.1
                : 500. 正答率(test)=0.1135
Generation: 600. 正答率(train)=0.13
                : 600. 正答率(test)=0.1135
Generation: 700. 正答率(train)=0.13
                : 700. 正答率(test)=0.1135
Generation: 800. 正答率(train)=0.12
                : 800. 正答率(test)=0.1135
Generation: 900. 正答率(train)=0.08
                : 900. 正答率(test)=0.1135
Generation: 1000. 正答率(train)=0.12
                : 1000. 正答率(test)=0.1135

```

```
In [26]: # グラフ表示
lists = range(0, iters_num, plot_interval)
plt.plot(lists, accuracies_train, label="training set")
plt.plot(lists, accuracies_test, label="test set")
plt.legend()
plt.title("accuracy: SGD")
plt.xlabel("iter")
plt.ylabel("accuracy")
plt.ylim(0, 1.0)
plt.show()
```



オフティマイザ : MOMENTUM

```
In [20]: use_batchnorm = False

network = MultiLayerNet(
    input_size=784, hidden_size_list=[40, 20], output_size=10, activation='sigmoid',
    weight_init_std=0.01, use_batchnorm=use_batchnorm)
```

```
In [21]: iters_num = 1000
train_size = x_train.shape[0]
batch_size = 100
learning_rate = 0.01
# 慣性 :  $\mu$ 
momentum = 0.9

plot_interval=10
```

```

In [22]: train_loss_list = []
         accuracies_train = []
         accuracies_test = []

         for i in range(iters_num):
             batch_mask = np.random.choice(train_size, batch_size)
             x_batch = x_train[batch_mask]
             d_batch = d_train[batch_mask]

             grad = network.gradient(x_batch, d_batch)

             if i == 0:
                 v = {}

             for key in ('W1', 'W2', 'W3', 'b1', 'b2', 'b3'):
                 if i == 0:
                     # イテレーションの最初に、vを初期化する
                     v[key] = np.zeros_like(network.params[key])

                 # momentumの更新量を計算
                 v[key] = momentum * v[key] - learning_rate * grad[key]
                 # パラメータの更新
                 network.params[key] += v[key]

                 # 誤差の計算
                 loss = network.loss(x_batch, d_batch)
                 train_loss_list.append(loss)

             if (i + 1) % plot_interval == 0:
                 accr_test = network.accuracy(x_test, d_test)
                 accuracies_test.append(accr_test)
                 accr_train = network.accuracy(x_batch, d_batch)
                 accuracies_train.append(accr_train)

                 if (i + 1) % (plot_interval * 10) == 0:
                     print(f'Generation: {i+1}. 正答率(train)={accr_train}')
                     print(f'          : {i+1}. 正答率(test)={accr_test}')


```

```

Generation: 100. 正答率(train)=0.08
          : 100. 正答率(test)=0.1135
Generation: 200. 正答率(train)=0.11
          : 200. 正答率(test)=0.1028
Generation: 300. 正答率(train)=0.09
          : 300. 正答率(test)=0.1135
Generation: 400. 正答率(train)=0.11
          : 400. 正答率(test)=0.1135
Generation: 500. 正答率(train)=0.14
          : 500. 正答率(test)=0.098
Generation: 600. 正答率(train)=0.07
          : 600. 正答率(test)=0.1135
Generation: 700. 正答率(train)=0.12
          : 700. 正答率(test)=0.1135
Generation: 800. 正答率(train)=0.26
          : 800. 正答率(test)=0.1135
Generation: 900. 正答率(train)=0.14
          : 900. 正答率(test)=0.098
Generation: 1000. 正答率(train)=0.09
          : 1000. 正答率(test)=0.1032

```

```
In [25]: lists = range(0, iters_num, plot_interval)
plt.plot(lists, accuracies_train, label="training set")
plt.plot(lists, accuracies_test, label="test set")
plt.legend()
plt.title("accuracy: MOMENTUM")
plt.xlabel("iter")
plt.ylabel("accuracy")
plt.ylim(0, 1.0)
plt.show()
```



MomentumをもとにAdaGradを作ってみよう

$\theta = 1e-4$ とする

```
In [48]: use_batchnorm = False

network = MultiLayerNet(
    input_size=784, hidden_size_list=[40, 20], output_size=10, activation='sigmoid',
    weight_init_std=0.01, use_batchnorm=use_batchnorm)
```

```
In [49]: iters_num = 1000

train_size = x_train.shape[0]
batch_size = 100
learning_rate = 0.01

# momentumは使用しない
theta = 1e-4

plot_interval=10
```

```

In [50]: train_loss_list = []
         accuracies_train = []
         accuracies_test = []

         for i in range(iters_num):
             batch_mask = np.random.choice(train_size, batch_size)
             x_batch = x_train[batch_mask]
             d_batch = d_train[batch_mask]

             grad = network.gradient(x_batch, d_batch)

             if i == 0:
                 h = {}

             for key in ('W1', 'W2', 'W3', 'b1', 'b2', 'b3'):
                 # 変更しよう =====
                 if i == 0:
                     h[key] = np.zeros_like(theta) # hの初期値はθ

                 h[key] = grad[key] * grad[key]
                 network.params[key] -= learning_rate * grad[key] / (np.sqrt(h[key]) + theta)
                 # =====

             loss = network.loss(x_batch, d_batch)
             train_loss_list.append(loss)

             if (i + 1) % plot_interval == 0:
                 accr_test = network.accuracy(x_test, d_test)
                 accuracies_test.append(accr_test)
                 accr_train = network.accuracy(x_batch, d_batch)
                 accuracies_train.append(accr_train)

                 if (i + 1) % (plot_interval * 10) == 0:
                     print(f'Generation: {i+1}. 正答率(train)={accr_train}')
                     print(f'          : {i+1}. 正答率(test)={accr_test}')


```

```

Generation: 100. 正答率(train)=0.26
                : 100. 正答率(test)=0.2645
Generation: 200. 正答率(train)=0.4
                : 200. 正答率(test)=0.3406
Generation: 300. 正答率(train)=0.61
                : 300. 正答率(test)=0.6339
Generation: 400. 正答率(train)=0.84
                : 400. 正答率(test)=0.8421
Generation: 500. 正答率(train)=0.89
                : 500. 正答率(test)=0.894
Generation: 600. 正答率(train)=0.94
                : 600. 正答率(test)=0.9115
Generation: 700. 正答率(train)=0.98
                : 700. 正答率(test)=0.9141
Generation: 800. 正答率(train)=0.97
                : 800. 正答率(test)=0.9236
Generation: 900. 正答率(train)=0.97
                : 900. 正答率(test)=0.9348
Generation: 1000. 正答率(train)=0.96
                : 1000. 正答率(test)=0.9309

```

```
In [51]: lists = range(0, iters_num, plot_interval)
plt.plot(lists, accuracies_train, label="training set")
plt.plot(lists, accuracies_test, label="test set")
plt.legend()
plt.title("accuracy: AdaGrad")
plt.xlabel("iter")
plt.ylabel("accuracy")
plt.ylim(0, 1.0)
plt.show()
```



オフティマイザ : RMSProp

```
In [31]: use_batchnorm = False

network = MultiLayerNet(
    input_size=784, hidden_size_list=[40, 20], output_size=10, activation='sigmoid',
    weight_init_std=0.01, use_batchnorm=use_batchnorm)
```

```
In [32]: iters_num = 1000
train_size = x_train.shape[0]
batch_size = 100
learning_rate = 0.01
decay_rate = 0.99

plot_interval=10
```

```

In [33]: train_loss_list = []
         accuracies_train = []
         accuracies_test = []

         for i in range(iters_num):
             batch_mask = np.random.choice(train_size, batch_size)
             x_batch = x_train[batch_mask]
             d_batch = d_train[batch_mask]

             grad = network.gradient(x_batch, d_batch)

             if i == 0:
                 h = {}

             for key in ('W1', 'W2', 'W3', 'b1', 'b2', 'b3'):
                 if i == 0:
                     h[key] = np.zeros_like(network.params[key])

                 h[key] *= decay_rate
                 h[key] += (1 - decay_rate) * np.square(grad[key])
                 network.params[key] -= learning_rate * grad[key] / (np.sqrt(h[key]) + 1e-7)

             loss = network.loss(x_batch, d_batch)
             train_loss_list.append(loss)

             if (i + 1) % plot_interval == 0:
                 accr_test = network.accuracy(x_test, d_test)
                 accuracies_test.append(accr_test)
                 accr_train = network.accuracy(x_batch, d_batch)
                 accuracies_train.append(accr_train)

                 if (i + 1) % (plot_interval * 10) == 0:
                     print(f'Generation: {i+1}. 正答率(train)={accr_train}')
                     print(f'          : {i+1}. 正答率(test)={accr_test}')

```


```

Generation: 100. 正答率(train)=0.66
          : 100. 正答率(test)=0.595
Generation: 200. 正答率(train)=0.94
          : 200. 正答率(test)=0.8714
Generation: 300. 正答率(train)=0.94
          : 300. 正答率(test)=0.9111
Generation: 400. 正答率(train)=0.96
          : 400. 正答率(test)=0.9102
Generation: 500. 正答率(train)=0.96
          : 500. 正答率(test)=0.9276
Generation: 600. 正答率(train)=1.0
          : 600. 正答率(test)=0.931
Generation: 700. 正答率(train)=0.96
          : 700. 正答率(test)=0.9338
Generation: 800. 正答率(train)=0.99
          : 800. 正答率(test)=0.9396
Generation: 900. 正答率(train)=0.99
          : 900. 正答率(test)=0.9395
Generation: 1000. 正答率(train)=0.97
          : 1000. 正答率(test)=0.931

```



```
In [34]: lists = range(0, iters_num, plot_interval)
plt.plot(lists, accuracies_train, label="training set")
plt.plot(lists, accuracies_test, label="test set")
plt.legend()
plt.title("accuracy: RMSProp")
plt.xlabel("iter")
plt.ylabel("accuracy")
plt.ylim(0, 1.0)
plt.show()
```



オフティマイザ : Adam

```
In [35]: use_batchnorm = False

network = MultiLayerNet(
    input_size=784, hidden_size_list=[40, 20], output_size=10, activation='sigmoid',
    weight_init_std=0.01, use_batchnorm=use_batchnorm)
```

```
In [36]: iters_num = 1000
train_size = x_train.shape[0]
batch_size = 100
learning_rate = 0.01
beta1 = 0.9
beta2 = 0.999

plot_interval=10
```

```

In [37]: train_loss_list = []
         accuracies_train = []
         accuracies_test = []

         for i in range(iters_num):
             batch_mask = np.random.choice(train_size, batch_size)
             x_batch = x_train[batch_mask]
             d_batch = d_train[batch_mask]

             grad = network.gradient(x_batch, d_batch)
             if i == 0:
                 m = {}
                 v = {}

             learning_rate_t = learning_rate * np.sqrt(1.0 - beta2 ** (i + 1)) / (1.0 - beta1 ** (i + 1))

             for key in ('W1', 'W2', 'W3', 'b1', 'b2', 'b3'):
                 if i == 0:
                     m[key] = np.zeros_like(network.params[key])
                     v[key] = np.zeros_like(network.params[key])

                 m[key] += (1 - beta1) * (grad[key] - m[key])
                 v[key] += (1 - beta2) * (grad[key] ** 2 - v[key])
                 network.params[key] -= learning_rate_t * m[key] / (np.sqrt(v[key]) + 1e-7)

             if (i + 1) % plot_interval == 0:
                 accr_test = network.accuracy(x_test, d_test)
                 accuracies_test.append(accr_test)
                 accr_train = network.accuracy(x_batch, d_batch)
                 accuracies_train.append(accr_train)
                 loss = network.loss(x_batch, d_batch)
                 train_loss_list.append(loss)

             if (i + 1) % (plot_interval * 10) == 0:
                 print(f'Generation: {i+1}. 正答率(train)={accr_train}')
                 print(f'          : {i+1}. 正答率(test)={accr_test}')

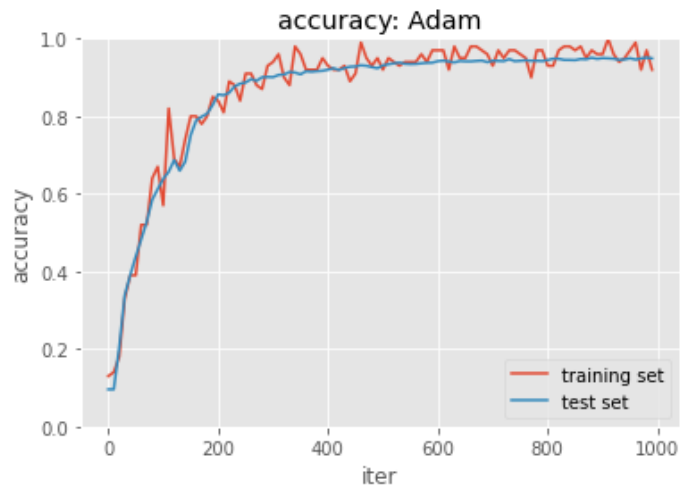
```

```

Generation: 100. 正答率(train)=0.67
          : 100. 正答率(test)=0.6119
Generation: 200. 正答率(train)=0.85
          : 200. 正答率(test)=0.829
Generation: 300. 正答率(train)=0.93
          : 300. 正答率(test)=0.9013
Generation: 400. 正答率(train)=0.95
          : 400. 正答率(test)=0.9172
Generation: 500. 正答率(train)=0.95
          : 500. 正答率(test)=0.9237
Generation: 600. 正答率(train)=0.97
          : 600. 正答率(test)=0.9381
Generation: 700. 正答率(train)=0.96
          : 700. 正答率(test)=0.9395
Generation: 800. 正答率(train)=0.97
          : 800. 正答率(test)=0.9423
Generation: 900. 正答率(train)=0.96
          : 900. 正答率(test)=0.948
Generation: 1000. 正答率(train)=0.92
          : 1000. 正答率(test)=0.9491

```

```
In [38]: lists = range(0, iters_num, plot_interval)
plt.plot(lists, accuracies_train, label="training set")
plt.plot(lists, accuracies_test, label="test set")
plt.legend()
plt.title("accuracy: Adam")
plt.xlabel("iter")
plt.ylabel("accuracy")
plt.ylim(0, 1.0)
plt.show()
```



考察：

SGDとMomentumは、イテレーションを進めても学習が進まない。

AdaGradは、iter=200くらいまでは精度の上昇が緩やかだが、それ以降は順調に学習が進んでいった。

RMSPropとAdamは、早い段階から学習が進んでいった。

In []: