

実装演習：2-3.過学習

```
In [24]: import numpy as np
from collections import OrderedDict
from common import layers
from data.mnist import load_mnist
import matplotlib.pyplot as plt
from multi_layer_net import MultiLayerNet
from common import optimizer
plt.style.use('ggplot')
```

```
In [2]: # MNISTデータ読み込み
(x_train, d_train), (x_test, d_test) = load_mnist(normalize=True)
```

```
In [3]: # 過学習を再現するために、学習データを削減
x_train = x_train[:300]
d_train = d_train[:300]
```

```
In [4]: # モデル構築
network = MultiLayerNet(input_size=784,
                        hidden_size_list=[100, 100, 100, 100, 100, 100],
                        output_size=10)

# オプティマイザ
optimizer = optimizer.SGD(learning_rate=0.01)
```

```
In [5]: iters_num = 1000
train_size = x_train.shape[0]
batch_size = 100

plot_interval=10
```

```

In [6]: train_loss_list = []
        accuracies_train = []
        accuracies_test = []

        for i in range(iters_num):
            batch_mask = np.random.choice(train_size, batch_size)
            x_batch = x_train[batch_mask]
            d_batch = d_train[batch_mask]

            grad = network.gradient(x_batch, d_batch)
            optimizer.update(network.params, grad)

            loss = network.loss(x_batch, d_batch)
            train_loss_list.append(loss)

            if (i+1) % plot_interval == 0:
                accr_train = network.accuracy(x_train, d_train)
                accr_test = network.accuracy(x_test, d_test)
                accuracies_train.append(accr_train)
                accuracies_test.append(accr_test)

                if (i+1) % (plot_interval * 10) == 0:
                    print(f'iter: {i+1}. accr_train={accr_train}, accr_test={accr_test}')

```

```


iter: 100. accr_train=0.7033333333333334, accr_test=0.5186
iter: 200. accr_train=0.9433333333333334, accr_test=0.7014
iter: 300. accr_train=0.9933333333333333, accr_test=0.7352
iter: 400. accr_train=1.0, accr_test=0.745
iter: 500. accr_train=1.0, accr_test=0.7486
iter: 600. accr_train=1.0, accr_test=0.7526
iter: 700. accr_train=1.0, accr_test=0.7567
iter: 800. accr_train=1.0, accr_test=0.7558
iter: 900. accr_train=1.0, accr_test=0.7566
iter: 1000. accr_train=1.0, accr_test=0.7571

```

```

In [7]: lists = range(0, iters_num, plot_interval)
        plt.plot(lists, accuracies_train, label="training set")
        plt.plot(lists, accuracies_test, label="test set")
        plt.legend()
        plt.title("accuracy")
        plt.xlabel("iter")
        plt.ylabel("accuracy")
        plt.ylim(0, 1.1)
        plt.show()

```



weight decay (重み減衰) による過学習の改善

L2正則化

```
In [8]: network = MultiLayerNet(input_size=784, hidden_size_list=[100, 100, 100, 100, 100, 100],  
output_size=10)
```

```
In [9]: iters_num = 1000  
train_size = x_train.shape[0]  
batch_size = 100  
learning_rate=0.01  
  
plot_interval=10  
hidden_layer_num = network.hidden_layer_num  
print(f'hidden_layer_num: {hidden_layer_num}')  
  
weight_decay_lambda = 0.1  
print(f'正則化強度設定: {weight_decay_lambda}')  
  
hidden_layer_num: 6  
正則化強度設定: 0.1
```

```

In [10]: train_loss_list = []
          accuracies_train = []
          accuracies_test = []

          for i in range(iters_num):
              batch_mask = np.random.choice(train_size, batch_size)
              x_batch = x_train[batch_mask]
              d_batch = d_train[batch_mask]

              grad = network.gradient(x_batch, d_batch)
              weight_decay = 0

              for idx in range(1, hidden_layer_num+1):
                  grad['W' + str(idx)] = network.layers['Affine' + str(idx)].dW + weight_decay_lambda *
network.params['W' + str(idx)]
                  grad['b' + str(idx)] = network.layers['Affine' + str(idx)].db

                  network.params['W' + str(idx)] -= learning_rate * grad['W' + str(idx)]
                  network.params['b' + str(idx)] -= learning_rate * grad['b' + str(idx)]

                  weight_decay += 0.5 * weight_decay_lambda * np.sqrt(np.sum(network.params['W'
+ str(idx)] ** 2))

              loss = network.loss(x_batch, d_batch) + weight_decay
              train_loss_list.append(loss)

              if (i+1) % plot_interval == 0:
                  accr_train = network.accuracy(x_train, d_train)
                  accr_test = network.accuracy(x_test, d_test)
                  accuracies_train.append(accr_train)
                  accuracies_test.append(accr_test)

              if (i+1) % (plot_interval * 10) == 0:
                  print(f'iter: {i+1}. accr_train={accr_train}, accr_test={accr_test}')


```

```

iter: 100. accr_train=0.61, accr_test=0.4161
iter: 200. accr_train=0.7566666666666667, accr_test=0.5982
iter: 300. accr_train=0.87, accr_test=0.6722
iter: 400. accr_train=0.88, accr_test=0.7053
iter: 500. accr_train=0.91, accr_test=0.7075
iter: 600. accr_train=0.9033333333333333, accr_test=0.7109
iter: 700. accr_train=0.91, accr_test=0.7119
iter: 800. accr_train=0.9066666666666666, accr_test=0.7043
iter: 900. accr_train=0.91, accr_test=0.7081
iter: 1000. accr_train=0.9133333333333333, accr_test=0.7099

```

```
In [11]: lists = range(0, iters_num, plot_interval)
plt.plot(lists, accuracies_train, label="training set")
plt.plot(lists, accuracies_test, label="test set")
plt.legend()
plt.title("accuracy: L2 Regularization")
plt.xlabel("iter")
plt.ylabel("accuracy")
plt.ylim(0, 1.1)
plt.show()
```



L1正則化

```
In [12]: network = MultiLayerNet(input_size=784, hidden_size_list=[100, 100, 100, 100, 100, 100],
output_size=10)
```

```
In [13]: iters_num = 1000
train_size = x_train.shape[0]
batch_size = 100
learning_rate=0.1

plot_interval=10
hidden_layer_num = network.hidden_layer_num
print(f'hidden_layer_num: {hidden_layer_num}')

weight_decay_lambda = 0.005
print(f'正則化強度設定: {weight_decay_lambda}')
```

```
hidden_layer_num: 6
正則化強度設定: 0.005
```

```

In [14]: train_loss_list = []
         accuracies_train = []
         accuracies_test = []

         for i in range(iters_num):
             batch_mask = np.random.choice(train_size, batch_size)
             x_batch = x_train[batch_mask]
             d_batch = d_train[batch_mask]

             grad = network.gradient(x_batch, d_batch)
             weight_decay = 0

             for idx in range(1, hidden_layer_num+1):
                 grad['W' + str(idx)] = network.layers['Affine' + str(idx)].dW + weight_decay_lambda *
np.sign(network.params['W' + str(idx)])
                 grad['b' + str(idx)] = network.layers['Affine' + str(idx)].db

                 network.params['W' + str(idx)] -= learning_rate * grad['W' + str(idx)]
                 network.params['b' + str(idx)] -= learning_rate * grad['b' + str(idx)]

                 weight_decay += weight_decay_lambda * np.sum(np.abs(network.params['W' + str(i
dx)]))

             loss = network.loss(x_batch, d_batch) + weight_decay
             train_loss_list.append(loss)

             if (i+1) % plot_interval == 0:
                 accr_train = network.accuracy(x_train, d_train)
                 accr_test = network.accuracy(x_test, d_test)
                 accuracies_train.append(accr_train)
                 accuracies_test.append(accr_test)

                 if (i+1) % (plot_interval * 10) == 0:
                     print(f'iter: {i+1}. accr_train={accr_train}, accr_test={accr_test}')


```

```

iter: 100. accr_train=0.91, accr_test=0.6895
iter: 200. accr_train=0.8366666666666667, accr_test=0.6201
iter: 300. accr_train=0.8333333333333334, accr_test=0.6451
iter: 400. accr_train=0.9, accr_test=0.6933
iter: 500. accr_train=0.63, accr_test=0.5113
iter: 600. accr_train=0.9633333333333334, accr_test=0.7625
iter: 700. accr_train=0.65, accr_test=0.4889
iter: 800. accr_train=0.9666666666666667, accr_test=0.7462
iter: 900. accr_train=0.9833333333333333, accr_test=0.7662
iter: 1000. accr_train=0.91, accr_test=0.6668

```

```
In [15]: lists = range(0, iters_num, plot_interval)
plt.plot(lists, accuracies_train, label="training set")
plt.plot(lists, accuracies_test, label="test set")
plt.legend()
plt.title("accuracy: L1 Regularization")
plt.xlabel("iter")
plt.ylabel("accuracy")
plt.ylim(0, 1.1)
plt.show()
```



Dropout

```
In [19]: class Dropout:
def __init__(self, dropout_ratio=0.5):
    self.dropout_ratio = dropout_ratio
    self.mask = None

def forward(self, x, train_flg=True):
    if train_flg:
        self.mask = np.random.rand(*x.shape) > self.dropout_ratio
        return x * self.mask
    else:
        return x * (1.0 - self.dropout_ratio)

def backward(self, dout):
    return dout * self.mask
```

```
In [20]: # ドロップアウト設定
use_dropout = True
dropout_ratio = 0.15
```

```
In [21]: network = MultiLayerNet(input_size=784, hidden_size_list=[100, 100, 100, 100, 100, 100],
output_size=10,
weight_decay_lambda=weight_decay_lambda, use_dropout=use_dropout,
dropout_ratio=dropout_ratio)
```

```
In [25]: sgd = optimizer.SGD(learning_rate=0.01)
```

```
In [26]: iters_num = 1000
train_size = x_train.shape[0]
batch_size = 100
plot_interval=10
```

```

In [27]: train_loss_list = []
        accuracies_train = []
        accuracies_test = []

        for i in range(iters_num):
            batch_mask = np.random.choice(train_size, batch_size)
            x_batch = x_train[batch_mask]
            d_batch = d_train[batch_mask]

            grad = network.gradient(x_batch, d_batch)
            sgd.update(network.params, grad)

            loss = network.loss(x_batch, d_batch)
            train_loss_list.append(loss)

            if (i+1) % plot_interval == 0:
                accr_train = network.accuracy(x_train, d_train)
                accr_test = network.accuracy(x_test, d_test)
                accuracies_train.append(accr_train)
                accuracies_test.append(accr_test)

                if (i+1) % (plot_interval * 10) == 0:
                    print(f'iter: {i+1}. accr_train={accr_train}, accr_test={accr_test}')

```

```


iter: 100. accr_train=0.3333333333333333, accr_test=0.2591
iter: 200. accr_train=0.48333333333333334, accr_test=0.383
iter: 300. accr_train=0.6266666666666667, accr_test=0.4921
iter: 400. accr_train=0.7066666666666667, accr_test=0.554
iter: 500. accr_train=0.7166666666666667, accr_test=0.5759
iter: 600. accr_train=0.7466666666666667, accr_test=0.6011
iter: 700. accr_train=0.8, accr_test=0.618
iter: 800. accr_train=0.84, accr_test=0.6512
iter: 900. accr_train=0.8666666666666667, accr_test=0.6662
iter: 1000. accr_train=0.92, accr_test=0.6877

```

```

In [28]: lists = range(0, iters_num, plot_interval)
        plt.plot(lists, accuracies_train, label="training set")
        plt.plot(lists, accuracies_test, label="test set")
        plt.legend()
        plt.title("accuracy: Dropout")
        plt.xlabel("iter")
        plt.ylabel("accuracy")
        plt.ylim(0, 1.1)
        plt.show()

```




```
In [29]: # ドロップアウト設定
use_dropout = True
dropout_ratio = 0.08
```

```
In [30]: network = MultiLayerNet(input_size=784, hidden_size_list=[100, 100, 100, 100, 100, 100],
                                output_size=10,
                                use_dropout = use_dropout, dropout_ratio = dropout_ratio)
```

```
In [31]: iters_num = 1000
train_size = x_train.shape[0]
batch_size = 100
learning_rate=0.01

hidden_layer_num = network.hidden_layer_num
plot_interval=10

# 正則化強度設定
weight_decay_lambda=0.004
```

```
In [32]: train_loss_list = []
accuracies_train = []
accuracies_test = []

for i in range(iters_num):
    batch_mask = np.random.choice(train_size, batch_size)
    x_batch = x_train[batch_mask]
    d_batch = d_train[batch_mask]

    grad = network.gradient(x_batch, d_batch)
    weight_decay = 0

    for idx in range(1, hidden_layer_num+1):
        grad['W' + str(idx)] = network.layers['Affine' + str(idx)].dW + weight_decay_lambda *
np.sign(network.params['W' + str(idx)])
        grad['b' + str(idx)] = network.layers['Affine' + str(idx)].db
        network.params['W' + str(idx)] -= learning_rate * grad['W' + str(idx)]
        network.params['b' + str(idx)] -= learning_rate * grad['b' + str(idx)]
        weight_decay += weight_decay_lambda * np.sum(np.abs(network.params['W' + str(i
dx)])))


    loss = network.loss(x_batch, d_batch) + weight_decay
    train_loss_list.append(loss)

    if (i+1) % plot_interval == 0:
        accr_train = network.accuracy(x_train, d_train)
        accr_test = network.accuracy(x_test, d_test)
        accuracies_train.append(accr_train)
        accuracies_test.append(accr_test)

    if (i+1) % (plot_interval * 10) == 0:
        print(f'iter: {i+1}. accr_train={accr_train}, accr_test={accr_test}')
```

```
iter: 100. accr_train=0.28, accr_test=0.2225
iter: 200. accr_train=0.43666666666666665, accr_test=0.3428
iter: 300. accr_train=0.49, accr_test=0.3866
iter: 400. accr_train=0.5966666666666667, accr_test=0.4741
iter: 500. accr_train=0.6933333333333334, accr_test=0.5883
iter: 600. accr_train=0.76, accr_test=0.6215
iter: 700. accr_train=0.82, accr_test=0.6475
iter: 800. accr_train=0.88, accr_test=0.6772
iter: 900. accr_train=0.8866666666666667, accr_test=0.6995
iter: 1000. accr_train=0.9166666666666666, accr_test=0.7115
```

```
In [33]: lists = range(0, iters_num, plot_interval)
plt.plot(lists, accuracies_train, label="training set")
plt.plot(lists, accuracies_test, label="test set")
plt.legend()
plt.title("accuracy: Dropout + L1 Regularization")
plt.xlabel("iter")
plt.ylabel("accuracy")
plt.ylim(0, 1.1)
plt.show()
```



考察：

最初の何もしないデータでは、学習を進めると訓練セットに対しては精度100%になるのに対し、テストセットに対して75%程度となった。これは訓練セット>テストセットで過学習と言える。また、精度が100%になる時点で、過学習を疑うべきである。

L2正則化では、過学習自体は改善していないものの、訓練セットの精度が91%程度で収束している。何もしないときに比べ、過学習は多少改善されているように見える。

L1正則化では、グラフがガクガクに変化している。これは、特徴量を一部完全に捨てていることに起因する。

ドロップアウトをすると、学習が収束するまで時間がかかるようになった。

ドロップアウト+L1正則化で、ドロップアウトしただけに比べ、テストセットの精度が68%から71%まで改善した。

In []: